

Supervised Learning : Regression

This notebook is the solution for the project under IBM Machine Learning Certificate. The contents of the notebook are listed below:

Exploration of data set

1. Exploratory Data Analytics
 - a. Data cleaning
 - b. Types of features
 - Numerical features
 - Categorical features
 - c. Statistical analysis of Data set
 - d. Skewness of numerical features
2. Split the Train test split

Variations in models

1. Linear regression model
 - One-hot encoding
 - Square root transformation
 - Standard scaling
 - Polynomial features
2. Cross validation
 - K-fold validation

Conclusion

▼ Import Header files

```
import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
import seaborn as sns
sns.set()
%matplotlib inline
import io
import sys
from google.colab import files
```

▼ Exploration of Dataset

Name of the dataset: [Used car Cost estimation Data \(Toyota\)](#)

Read data

```
uploaded = files.upload()
df2 = pd.read_csv(io.BytesIO(uploaded['toyota.csv']))
```

No file chosen Upload widget is only available when the cell has been executed in the current browser session. Please rerun this cell to enable.
Saving toyota.csv to toyota.csv

Size of data set

```
df2.shape
```

```
(6738, 9)
```

a. Data Cleaning

Remove duplicate entries and check the size of the data set.

```
df2 = df2.drop_duplicates(keep='first').reset_index(drop=True)
df2.shape
```

(6699, 9)

There are 39 duplicate entries

b. Types of features

```
categorical_features = (df2.dtypes == np.object)
numerical_features = (df2.dtypes == np.float64) | (df2.dtypes == np.int64)

categorical_columns = df2.columns[categorical_features].tolist()
numerical_columns = df2.columns[numerical_features].tolist()

print(f'Categorical Feature columns: {categorical_columns}')
print(f'Numerical Feature columns: {numerical_columns}')

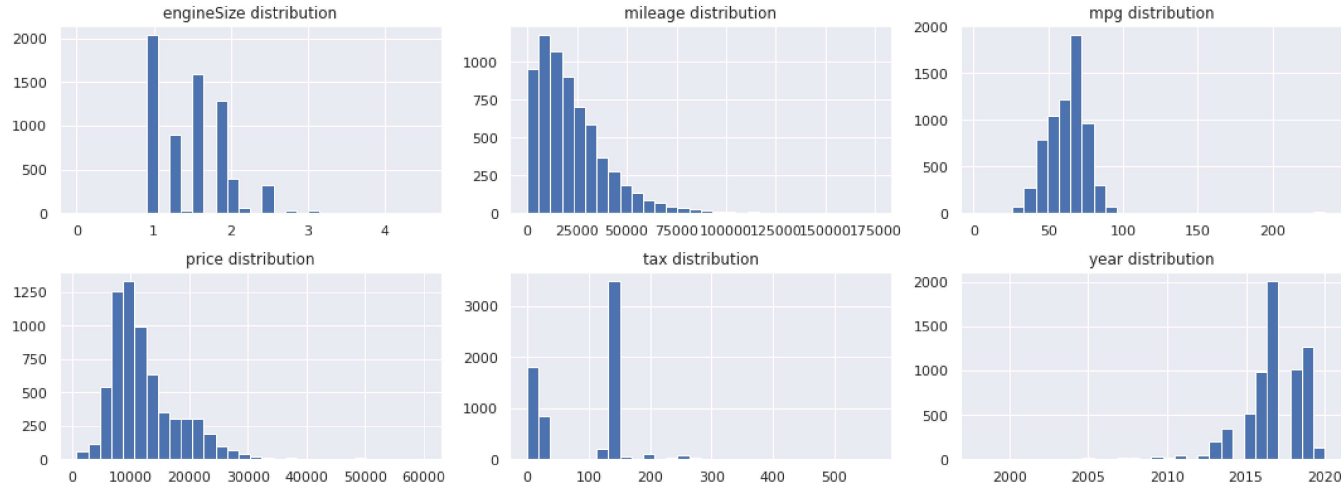
Categorical Feature columns: ['model', 'transmission', 'fuelType']
Numerical Feature columns: ['year', 'price', 'mileage', 'tax', 'mpg', 'engineSize']
```

c. Statistical Analysis of Data set

Plot histogram for numerical features

```
def hist_loop(input: pd.DataFrame,
              no_rows: int,
              no_cols: int,
              figure_size: tuple):
    fig, axes = plt.subplots(no_rows, no_cols, figsize=figure_size)
    for i, ax in enumerate(axes.flatten()):
        if i < len(input.columns):
            input[sorted(input.columns)[i]].plot.hist(bins=30, ax=ax)
            ax.set_title(f'{sorted(input.columns)[i]} distribution')
            ax.tick_params(axis='x')
            ax.tick_params(axis='y')
            ax.get_yaxis().get_label().set_visible(False)
        else:
            fig.delaxes(ax=ax)
    fig.tight_layout()
```

```
hist_loop(input=df2[numerical_columns],
          no_rows=3,
          no_cols=3,
          figure_size=(15,8))
```



d. Plot Skewness of numerical features

```
def function_skewness(input: pd.DataFrame, skewness_limit: float) -> pd.DataFrame:

    skewness_values = input.skew()
    skewness_columns = (skewness_values
                        .sort_values(ascending=False)
                        .to_frame('Skewness')
                        .query('abs(Skewness) > {}'.format(skewness_limit))
    )
    return skewness_columns
skewness_columns = function_skewness(df2[numerical_columns], 0.75)
```


skewness_columns

	Skewness
mpg	3.481386
mileage	1.910231
price	1.812441
engineSize	0.834013
year	-2.214188

Impression: Except tax distribution, other features are right skewed. In this case, square root transformation can be the best choice for it. Because tax distribution has zero entries.

```
input_sqrt = df2[numerical_columns].drop('price', axis=1).copy()
for cols in list(skewness_columns.index):
    if cols != 'price':
        input_sqrt[cols] = input_sqrt[cols].apply(np.sqrt)

# Check again
function_skewness(input_sqrt, 0.75)
sns.pairplot(input_sqrt.join(df2['price']), plot_kws=dict(alpha=.2, edgecolor='Red'));
```



Analysis of categorical features

```
df2.describe(include=np.object)
```

	model	transmission	fuelType
count	6699	6699	6699
unique	18	4	4
top	Yaris	Manual	Petrol
freq	2117	3793	4058

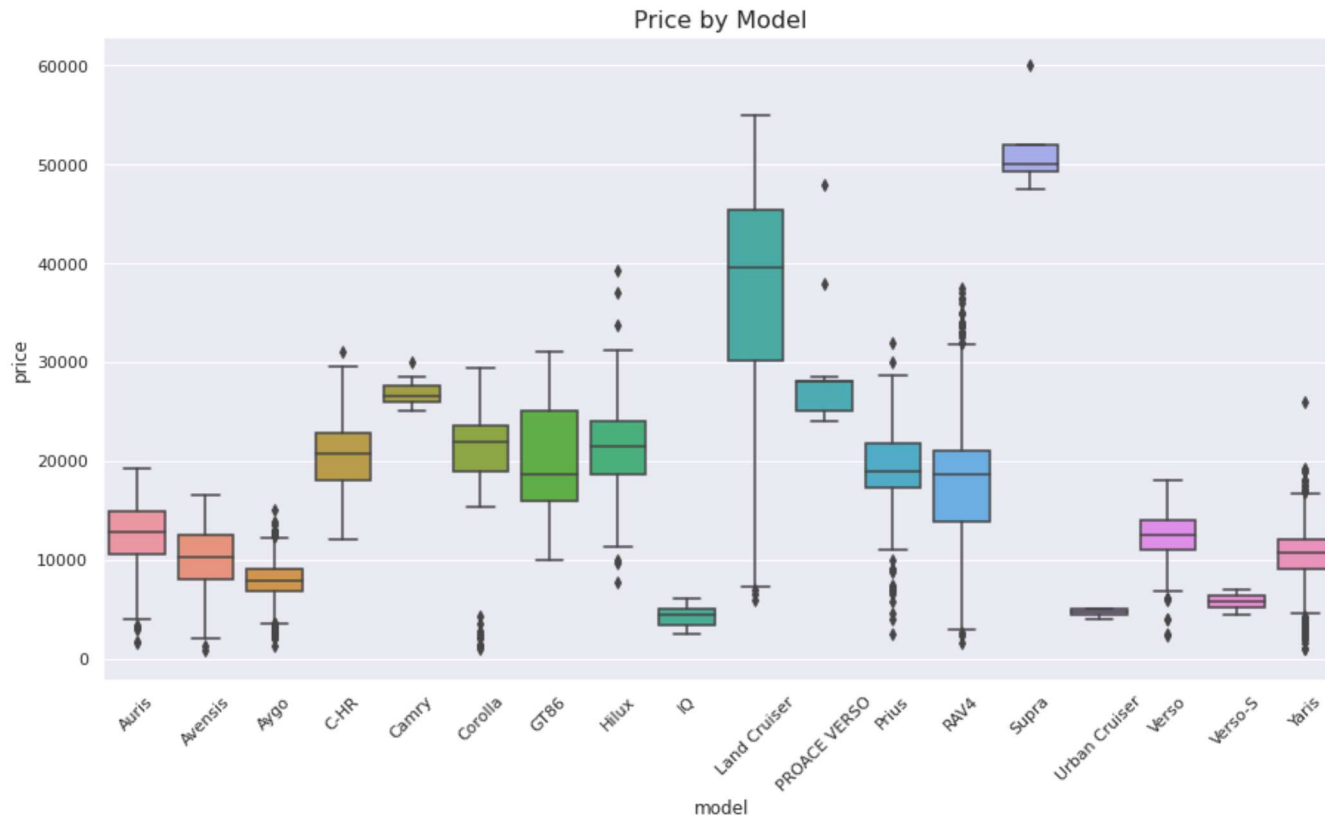
```
list(df2['model'].unique())
```

```
[' GT86',  
 ' Corolla',  
 ' RAV4',  
 ' Yaris',  
 ' Auris',  
 ' Aygo',  
 ' C-HR',  
 ' Prius',  
 ' Avensis',  
 ' Verso',  
 ' Hilux',  
 ' PROACE VERSO',  
 ' Land Cruiser',  
 ' Supra',  
 ' Camry',  
 ' Verso-S',  
 ' IQ',  
 ' Urban Cruiser']
```

```
df2['model'] = df2['model'].str.strip(' ')
```

```
fig, ax = plt.subplots(figsize=(15,8))  
order = sorted(list(df2['model'].unique()))  
sns.boxplot(x='model', v='price', data=df2, order=order, ax=ax)
```

```
plt.xticks(rotation=45)
plt.title('Price by Model', fontsize=16)
plt.show()
```



```
list(df2['transmission'].unique())
```

```
['Manual', 'Automatic', 'Semi-Auto', 'Other']
```


```
list(df2['fuelType'].unique())
```



```
['Petrol', 'Other', 'Hybrid', 'Diesel']
```

```
fig, ax = plt.subplots(1, 2, figsize=(20,6), sharey=True)
order0 = sorted(list(df2['transmission'].unique()))
sns.boxplot(x='transmission', y='price', data=df2, order=order0, ax=ax[0])
ax[0].set_title('Price by Transmission', fontsize=16)
ax[0].tick_params('x', labelrotation=45)

order1 = sorted(list(df2['fuelType'].unique()))
sns.boxplot(x='fuelType', y='price', data=df2, order=order1, ax=ax[1])
ax[1].set_title('Price by Fuel Type', fontsize=16)
ax[1].tick_params('x', labelrotation=45)
plt.show()
```



2. Test set split

Take 30% of the data set as test set.

```
from sklearn.model_selection import train_test_split
train_set, test_set = train_test_split(df2, test_size=0.3, random_state=0)
print(f'Training set size: {train_set.shape}')
print(f'Test set size: {test_set.shape}')
```

```
Training set size: (4689, 9)
Test set size: (2010, 9)
```

▼ Model variations

```
df = train_set.copy()
```

```
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error, r2_score
from sklearn.preprocessing import OneHotEncoder, StandardScaler, PolynomialFeatures
```

```
class LR_model:

    # Default values
    target = 'price'
    test_size = 0.25
    random_state = 0
    skew_cols = ['engineSize', 'mileage', 'age']

    def __init__(self, data):
        self.train, self.test = train_test_split(data, test_size=LR_model.test_size, random_state=LR_model.random_state)

    def clean_data(self, df):

        """ Cleans the data """

        df = df[df['year'] <= 2020]
        df['age'] = 2020 - df['year']
```

```
df = df.drop(['year'], axis=1)
df['model'] = df['model'].str.strip(' ')
```

```
return df
```

```
def oh_enc(self, X_train, X_test):
```

```
    """ Performs one-hot encoding and drops the first category
    """
```

```
    ENC = OneHotEncoder(handle_unknown='ignore', sparse=False)
```

```
    # Filter categorical features only
```

```
    X_train_cat = X_train.select_dtypes(include=['object'])
```

```
    X_test_cat = X_test.select_dtypes(include=['object'])
```

```
    # Fit one-hot encoding on training set
```

```
    # Transform both training set and test set
```

```
    X_train_enc = ENC.fit_transform(X_train_cat)
```

```
    X_test_enc = ENC.transform(X_test_cat)
```

```
    # Join dummy values with numerical features
```

```
    X_train_enc_df = pd.DataFrame(X_train_enc,
                                  index=X_train.index,
                                  columns=ENC.get_feature_names(X_train_cat.columns.tolist()))
```

```
    X_train = X_train_enc_df.join(X_train.select_dtypes(exclude=['object']))
```

```
    # Drop one column of each category
```

```
    for col in X_train_cat.columns.tolist():
```

```
        cat_cols = X_train.columns[X_train.columns.str.startswith(col)].tolist()
```

```
        if len(cat_cols) > 1:
```

```
            X_train = X_train.drop(cat_cols[0], axis=1)
```

```
    # Same steps for the test set
```

```
    X_test_enc_df = pd.DataFrame(X_test_enc,
                                  index=X_test.index,
                                  columns=ENC.get_feature_names(X_test_cat.columns.tolist()))
```

```
    X_test = X_test_enc_df.join(X_test.select_dtypes(exclude=['object']))
```

```
    for col in X_test_cat.columns.tolist():
```

```
        cat_cols = X_test.columns[X_test.columns.str.startswith(col)].tolist()
```

```
        if len(cat_cols) > 1:
```

```
            X_test = X_test.drop(cat_cols[0], axis=1)
```

```

X_test = X_test.drop(cat_cols[0], axis=1)

return X_train, X_test

def sqrt_trans(self, X_train, X_test):

    """ Applies square root transformation for skewed features
    """

    X_train[LR_model.skew_cols] = X_train[LR_model.skew_cols].apply(np.sqrt)
    X_test[LR_model.skew_cols] = X_test[LR_model.skew_cols].apply(np.sqrt)

    return X_train, X_test

def scale_X(self, X_train, X_test):

    """ Applies standard scaling for all numerical features
    """

    scaler = StandardScaler()

    # Filter numerical features only (excluding binary values)
    float_cols = X_train.columns[~X_train.isin([0,1]).all()].tolist()

    # Fit features in training set and transform to test set
    X_train[float_cols] = scaler.fit_transform(X_train[float_cols])
    X_test[float_cols] = scaler.transform(X_test[float_cols])

    return X_train, X_test

def add_pf(self, X_train, X_test, degree=None):

    """ Adds polynomial features into the dataset
    """

    PF = PolynomialFeatures(degree=degree, include_bias=False)

    # Filter numerical features only (excluding binary values)
    float_cols = X_train.columns[~X_train.isin([0,1]).all()].tolist()

    # Fit features in training set and transform to test set
    X_train_pf = PF.fit_transform(X_train[float_cols])
    X_test_pf = PF.transform(X_test[float_cols])

```

```

X_test_pf = PF.transform(X_test[~float_cols])

# Add non-numerical features back into the transformed training set
X_train_pf_df = pd.DataFrame(X_train_pf,
                              index=X_train.index,
                              columns=PF.get_feature_names(input_features=float_cols))
X_train = X_train_pf_df.join(X_train[X_train.columns[~X_train.columns.isin(float_cols)].tolist()])

# Add non-numerical features back into the transformed test set
X_test_pf_df = pd.DataFrame(X_test_pf,
                              index=X_test.index,
                              columns=PF.get_feature_names(input_features=float_cols))
X_test = X_test_pf_df.join(X_test[X_test.columns[~X_test.columns.isin(float_cols)].tolist()])

return X_train, X_test

def rmse(self, y_true, y_predicted):

    """ Returns root mean squared error"""
    return np.sqrt(mean_squared_error(y_true, y_predicted))

def model_fit(self, label=None,
              encoding=False,
              squareroot=False,
              scaling=False,
              polynomial=False,
              degree=None):

    """ Fits linear regression model and returns
    RMSE of training set and test set
    """

    # Split data
    train, test = self.clean_data(self.train), self.clean_data(self.test)
    X_train, X_test = train.drop(LR_model.target, axis=1), test.drop(LR_model.target, axis=1)
    y_train, y_test = train[LR_model.target], test[LR_model.target]

    # Perform feature engineering
    if encoding:
        X_train, X_test = self.oh_enc(X_train, X_test)

    if not encoding:

```

```

X_train, X_test = X_train.select_dtypes(exclude=['object']), X_test.select_dtypes(exclude=['object'])

if squareroot:
    X_train, X_test = self.sqrtrans(X_train, X_test)

if scaling:
    X_train, X_test = self.scale_X(X_train, X_test)

if polynomial:
    X_train, X_test = self.add_pf(X_train, X_test, degree)

num_features = X_train.shape[1]
# Fit model and predict the target
LR = LinearRegression()
LR.fit(X_train, y_train)
y_train_pred = LR.predict(X_train)
y_test_pred = LR.predict(X_test)

# Compute RMSE and store in a dictionary
rmse_train = self.rmse(y_train, y_train_pred)
rmse_test = self.rmse(y_test, y_test_pred)
scores = {
    'Model': label,
    'Number of features': num_features,
    'RMSE train': rmse_train,
    'RMSE test': rmse_test
}

return scores

```

One hot encoding

```

error_df = [] # Blank error list to create a data frame later

# Fit non-encoded data
not_enc = LR_model(df).model_fit(label='not encoded')

# Fit encoded data
enc = LR_model(df).model_fit(label='one hot encoded', encoding=True)

```

```
# Print out error table
error_df = pd.DataFrame.from_dict([not_enc])
error_df = error_df.append(enc, ignore_index=True)
error_df
```

	Model	Number of features	RMSE train	RMSE test
0	not encoded	5	3134.414974	3251.084281
1	one hot encoded	28	1700.386931	1960.896246

Square root transformation

```
not_enc_bc = LR_model(df).model_fit(label='not encoded + squareroot', squareroot=True)

# Fit encoded data
enc_bc = LR_model(df).model_fit(label='one hot encoded + squareroot', encoding=True, squareroot=True)

# Print out error table
error_df = error_df.append([not_enc_bc, enc_bc], ignore_index=True)
error_df
```

	Model	Number of features	RMSE train	RMSE test
0	not encoded	5	3134.414974	3251.084281
1	one hot encoded	28	1700.386931	1960.896246
2	not encoded + squareroot	5	3277.562670	3342.373357
3	one hot encoded + squareroot	28	1611.152344	1938.144640

Standard Scaling

```
enc_bc_s = LR_model(df).model_fit(label='one hot encoded + squareroot + scaled', encoding=True, squareroot=True, scaling=True)

# Print out error table
error_df = error_df.append(enc_bc_s, ignore_index=True)
error_df
```

error_df

	Model	Number of features	RMSE train	RMSE test
0	not encoded	5	3134.414974	3251.084281
1	one hot encoded	28	1700.386931	1960.896246
2	not encoded + squareroot	5	3277.562670	3342.373357
3	one hot encoded + squareroot	28	1611.152344	1938.144640
4	one hot encoded + squareroot + scaled	28	1611.152344	1938.144640

Polynomial Features

```
error_pf_df = pd.DataFrame(columns=['Model', 'Number of features', 'RMSE train', 'RMSE test'])

# Iterate different degree, 1 to 10
for d in list(range(1,11)):
    error = LR_model(df).model_fit(label=f'Degree = {d}', encoding=True, scaling=True, polynomial=True, degree=d)
    error_pf_df = error_pf_df.append(error, ignore_index=True)

# Print out the error table
error_pf_df
```


	Model	Number of features	RMSE train	RMSE test
0	Degree = 1	28	1700.386931	1.960896e+03

▼ Cross-validation and Regularization

3 Degree = 4 148 1157.390710 2.252540e+03

```
from sklearn.linear_model import Lasso, Ridge, ElasticNet
from sklearn.model_selection import KFold, cross_val_score
from sklearn.preprocessing import FunctionTransformer
from sklearn.pipeline import make_pipeline
from sklearn_pandas import DataFrameMapper, gen_features
```

1 Degree = 8 1309 678.600367 1.349548e+12

```
kf = KFold(shuffle=True, random_state=0, n_splits=5)
df = df[df['year'] <= 2020]
df['age'] = 2020 - df['year']
df = df.drop(['year'], axis=1)
df['model'] = df['model'].str.strip(' ')
```

```
X = df.drop('price', axis=1)
y = df['price']
```

```
class XPipe:
    # Skew features observed from the EDA
    skew_cols = ['engineSize', 'mileage', 'age']

    def __init__(self, X):
        self.cat_cols = X.select_dtypes(include=['object']).columns.tolist()
        self.num_cols = X.select_dtypes(exclude=['object']).columns.tolist()

    def drop_first(self, cat_values):

        """ Drops one category after one-hot encoding"""

        dummy_df = pd.DataFrame(cat_values)
        dummy_df = dummy_df.drop(0, axis=1)

        return dummy_df.values
```

```

def sqrt_trans(self, num_values):

    """ Applies square root transformation to skewed features"""

    num_df = pd.DataFrame(num_values, columns=self.num_cols)
    num_df[XPipe.skew_cols] = num_df[XPipe.skew_cols].apply(np.sqrt)

    return num_df.values

def model_pipe(self, model, degree=2):

    """ Returns a data pipeline"""

    cat_list = [[col] for col in self.cat_cols]
    cat_features = gen_features(
        columns=cat_list,
        classes=[{'class':OneHotEncoder, 'handle_unknown':'ignore', 'sparse':False},
                 {'class':FunctionTransformer, 'func': self.drop_first}]
    )

    mapper = DataFrameMapper((cat_features) + [
        (self.num_cols, [FunctionTransformer(self.sqrt_trans),
                          StandardScaler(),
                          PolynomialFeatures(degree=degree, include_bias=False)])
    ])

    return make_pipeline(mapper, model)

XPipe(X).model_pipe(LinearRegression())
metrics = {}

```

```

lr_scores = {}

# Iterate over different degrees, 1 to 6
for d in list(range(1,7)):
    score = cross_val_score(XPipe(X)
                            .model_pipe(LinearRegression(), degree=d),
                            X, y, cv=kf, scoring='neg_mean_squared_error')
    lr_scores[f'Degree = {d}'] = np.mean(np.sqrt(-score))

```

```
# Create an error table and print it out
lr_rmse = pd.DataFrame.from_dict(lr_scores, orient='index', columns=['Average RMSE']).sort_values('Average RMSE')
lr_rmse
```

	Average RMSE
Degree = 3	1432.183484
Degree = 2	1443.792111
Degree = 1	1718.408899
Degree = 4	1778.965184
Degree = 5	27755.681679
Degree = 6	180596.605059

```
lr_r2_scores = cross_val_score(XPipe(X)
                                .model_pipe(LinearRegression(), degree=3),
                                X, y, cv=kf)
lr_r2_scores
```

```
array([0.9630676 , 0.94695759, 0.96153676, 0.92694877, 0.95192922])
```

```
metrics['Average RMSE'] = [lr_rmse['Average RMSE'].values[0]]
metrics['Average R2'] = [np.mean(lr_r2_scores)]
```

▼ Lasso Regression

```
las_scores = {}

# Iterate over different alphas
for d in [1, 2, 3]:
    for a in [0.005, 0.01, 0.05, 0.1, 0.3, 1, 3, 5, 10]:
        score = cross_val_score(XPipe(X)
                                .model_pipe(Lasso(alpha=a, max_iter=100000), degree=d),
                                X, y, cv=kf, scoring='neg_mean_squared_error')
```

```

las_scores[f'Degree = {d}, alpha = {a}'] = np.mean(np.sqrt(-score))

# Create an error table and print it out (first 5 rows)
las_rmse = pd.DataFrame.from_dict(las_scores, orient='index', columns=['Average RMSE']).sort_values('Average RMSE')
las_rmse.head(5)

```

	Average RMSE
Degree = 3, alpha = 3	1393.401615
Degree = 3, alpha = 1	1399.678825
Degree = 3, alpha = 5	1413.715981
Degree = 3, alpha = 0.3	1418.124204
Degree = 3, alpha = 0.1	1427.361948

```

las_r2_scores = cross_val_score(XPipe(X)
                                .model_pipe(Lasso(alpha=0.3, max_iter=100000), degree=3),
                                X, y, cv=kf)

las_r2_scores

array([0.96331517, 0.94720766, 0.96139386, 0.93154873, 0.95265881])

```

```

metrics['Average RMSE'].append(las_rmse['Average RMSE'].values[0])
metrics['Average R2'].append(np.mean(las_r2_scores))

```

▼ Ridge Regression

```

ridge_scores = {}

# Iterate over different alphas
for d in [1, 2, 3]:
    for a in [0.005, 0.01, 0.05, 0.1, 0.3, 1, 3, 5, 10]:
        score = cross_val_score(XPipe(X)
                                .model_pipe(Ridge(alpha=a, max_iter=100000), degree=d),
                                X, y, cv=kf, scoring='neg_mean_squared_error')

```

```

        ridge_scores[f'Degree = {d}, alpha = {a}'] = np.mean(np.sqrt(-score))

# Create an error table and print it out (first 5 rows)
ridge_rmse = pd.DataFrame.from_dict(ridge_scores, orient='index', columns=['Average RMSE']).sort_values('Average RMSE')
ridge_rmse.head(5)

```

	Average RMSE
Degree = 3, alpha = 1	1398.916770
Degree = 3, alpha = 0.3	1403.848928
Degree = 3, alpha = 0.1	1418.622799
Degree = 3, alpha = 0.05	1424.711608
Degree = 3, alpha = 3	1426.107998

```

ridge_r2_scores = cross_val_score(XPipe(X)
                                .model_pipe(Ridge(alpha=0.005, max_iter=100000), degree=3),
                                X, y, cv=kf)

ridge_r2_scores

array([0.96309438, 0.94693861, 0.96153351, 0.92726481, 0.95194661])

```

```

metrics['Average RMSE'].append(ridge_rmse['Average RMSE'].values[0])
metrics['Average R2'].append(np.mean(ridge_r2_scores))

```

▼ Elastic Net Regression

```

elasticnet_scores = {}

# Iterate over different alphas
for d in [1, 2, 3]:
    for a in [0.005, 0.01, 0.05, 0.1, 0.3, 1, 3, 5, 10]:
        score = cross_val_score(XPipe(X)

```

```

score = cross_val_score(XPipe(X)
                          .model_pipe(ElasticNet(alpha=a, max_iter=100000), degree=d),
                          X, y, cv=kf, scoring='neg_mean_squared_error')
elasticnet_scores[f'Degree = {d}, alpha = {a}'] = np.mean(np.sqrt(-score))

# Create an error table and print it out (first 5 rows)
elasticnet_rmse = pd.DataFrame.from_dict(elasticnet_scores, orient='index', columns=['Average RMSE']).sort_values('Average RMSE')
elasticnet_rmse.head(5)

```

	Average RMSE
Degree = 3, alpha = 0.005	1471.017819
Degree = 3, alpha = 0.01	1510.381233
Degree = 2, alpha = 0.005	1583.090097
Degree = 2, alpha = 0.01	1638.951385
Degree = 3, alpha = 0.05	1693.665749

```

elasticnet_r2_scores = cross_val_score(XPipe(X)
                                      .model_pipe(ElasticNet(alpha=0.005, max_iter=100000), degree=3),
                                      X, y, cv=kf)

elasticnet_r2_scores

array([0.95893972, 0.94234671, 0.95363051, 0.93137833, 0.95321672])

```

```

metrics['Average RMSE'].append(elasticnet_rmse['Average RMSE'].values[0])
metrics['Average R2'].append(np.mean(elasticnet_r2_scores))

```

▼ Analysis of metrics

```

metrics['Model'] = ['Linear', 'Lasso', 'Ridge', 'Elastic Net']
pd.DataFrame.from_dict(metrics).set_index('Model').sort_values('Average RMSE')

```

	Average RMSE	Average R2
Model		
Lasso	1393.401615	0.951225
Ridge	1398.916770	0.950156
Linear	1432.183484	0.950088

Conclusion

Elastic Net registers the highest RMSE and Lasso model registers the maximum average R2

+ Code

+ Text